


# Motion Planning II - Sampling-based Planning

---

A lot of Material from Howie Choset, Nancy Amato, Sujay  
Bhattacharjee, G.D. Hager, S. LaValle, J. Kuffner, D. Hsu

# Last time...



- We learned about configuration space (C-space)



- How do we plan in **high-dimensional** C-spaces?


# Exact methods: The problem

- Exact methods either find a solution or prove none exists
- Require computing C-space obstacles
  - Very computationally expensive!
- Decomposition methods (cells, octree, etc.) also not practical b/c sweeping/decomposition is sensitive to dimension
  - Also requires either knowing C-space obstacle or lots of collision checking



# Discrete Planning: The problem

- Discrete search run-time and memory requirements are very sensitive to branching factor (number of successors)
- Number of successors depend on dimension
- For a 3-dimensional 8-connected space, 26 successors
- For an  $n$ -dimensional 8-connected space,  $3^n - 1$  successors
  - Increases very quickly!




# The problem

- Need a path planning method that isn't so sensitive to dimensionality
- But:
  - Path planning is PSPACE-hard  
[Reif 79, Hopcroft et al. 84, 86]
  - Complexity is exponential in dimension of the C-space [Canny 86]
- What if we weaken *completeness* and *optimality* requirements?



Real robots can have 20+ DOF!



# Weakening requirements

<u>Ideal</u>		<u>Practical in High Dimensions</u>
Complete	————→	Probabilistically Complete
Optimal	————→	Feasible

\*More recent methods show asymptotic optimality

- **Probabilistic completeness:** A path planner is probabilistically complete if, given a solvable problem, the probability that the planner solves the problem goes to 1 as time goes to infinity.
- Feasibility: Path obeys all constraints (usually obstacles).
- A feasible path can be optimized *locally* after it is found

# Sampling-based planning

- Main idea: Instead of systematically-discretizing the C-space, take samples in the C-space and use them to construct a graph




# Sampling-based planning

## Advantages

- Don't need to discretize C-space
- Don't need to explicitly represent C-space
- Easy to sample high-dimensional spaces

## Disadvantages

- Probability of sampling an area depends on the area's size
  - Hard to sample *narrow passages*
- No strict completeness/optimalty






# Outline

- PRM
- Sampling strategies
- RRT

# Probabilistic Roadmap (PRM)


- Main idea: Build a roadmap of the space from sampled points, search the roadmap to find a path
- Roadmap should capture the *connectivity* of the free space



# Probabilistic Roadmap (PRM)

- Building a PRM: 2 phase process
- “Learning” Phase (this is not really machine learning)
  - Construction Step
  - Expansion Step (not used in practice today)
- Query Phase
  - Answer a given path planning query
- PRMs are known as *multi-query algorithms*, because roadmap can be re-used if environment and robot haven’t changed between queries.

# PRM Example



## “Learning” Phase

- Construction step: Build the roadmap by sampling random free configurations and connect them using a fast *local planner*
- Store these configurations as nodes in a graph
  - Note: In PRM literature, nodes are sometimes called “milestones”
- Edges of the graph are the paths between nodes found by the local planner

# “Learning” Phase: Construction Step

Start with an empty graph  $G = (V, E)$

For  $i = 1$  to MaxIterations

Generate random configuration  $q$  ←

If  $q$  is collision-free

Add  $q$  to  $V$

Select  $k$  nearest nodes to  $q$  in  $V$  ←


Attempt connection between each of these nodes and  $q$  using local planner

If a connection is successful, add it as an edge in  $E$

# “Learning” Phase: Construction Step


Start with an empty graph  $G = (V, E)$

For  $i = 1$  to MaxIterations

Generate random configuration  $q$  

If  $q$  is collision-free

    Add  $q$  to  $V$

    Select  $k$  nearest nodes to  $q$  in  $V$  

    Attempt connection between each of these nodes and  $q$  using local planner

    If a connection is successful, add it as an edge in  $E$

# “Learning” Phase: Sampling Collision-free Configurations


- Easiest and most common: uniform random sampling in C-space
  - Draw random value in allowable range for each DOF, combine into a vector
  - Place robot at the configuration and check collision
  - Repeat above until you get a collision-free configuration
  - AKA “Rejection Sampling”
- MANY ways to do this, many papers published, we will discuss more methods later



# “Learning” Phase: Construction Step

Start with an empty graph  $G = (V, E)$

For  $i = 1$  to MaxIterations

Generate random configuration  $q$  

If  $q$  is collision-free

Add  $q$  to  $V$

Select  $k$  nearest nodes to  $q$  in  $V$  

Attempt connection between each of these nodes and  $q$  using local planner

If a connection is successful, add it as an edge in  $E$

## “Learning” Phase: Finding Nearest Neighbors (NN)


- Need to decide a distance metric  $D(q_1, q_2)$  to define “nearest”
- $D$  should reflect likelihood of success of local planner connection (roughly)
  - If  $D(q_1, q_2)$  is small, success should be likely
  - If  $D(q_1, q_2)$  is large, success should be less likely
- By default, use Euclidian distance:

$$D(q_1, q_2) = \|q_1 - q_2\|$$

- Can weigh different dimensions of C-space differently
  - Often used to weigh translation vs. rotation

# “Learning” Phase: Finding Nearest Neighbors (NN)

- Two popular ways to do NN in PRM
  - Find  $k$  nearest neighbors (even if they are distant)
  - Find all nearest neighbors within a certain distance
- Naïve NN computation can be slow with 1000s of nodes, so use *kd-tree* to store nodes and do NN queries
  - A *kd-tree* is a data-structure that recursively divides the space into bins that contain points (like Oct-tree and Quad-tree)
  - NN then searches through bins (not individual points) to find nearest point
  - Much faster to use *kd-tree* for large numbers of nodes
  - BUT, cost of constructing a *kd-tree* is significant, so only regenerate tree once in a while (not for every new node!)
  - *kd-tree* code is easy to find online



# “Learning” Phase: Construction Step

Start with an empty graph  $G = (V, E)$

For  $i = 1$  to MaxIterations

Generate random configuration  $q$  ←

If  $q$  is collision-free

    Add  $q$  to  $V$

    Select  $k$  nearest nodes to  $q$  in  $V$  ←

    Attempt connection between each of these nodes and  $q$  using local planner


    If a connection is successful, add it as an edge in  $E$

## “Learning” Phase: Local Planner

- In general, local planner can be anything that attempts to find a path between points, even another PRM!
- BUT, local planner needs to be fast b/c it's called many times by the algorithm
- Easiest and most common: Connect the two configurations with a straight line in C-space, check that the line is collision-free
  - Advantages:
    - Fast
    - Don't need to store local paths

## “Learning” Phase: Expansion step (not used in practice today)


- Problem: Can have disconnected components that should be connected
  - I.e. you haven't captured the true connectivity of the space



- Expansion step uses heuristics to sample more nodes in an effort to connect disconnected components
  - Unclear how to do this the “right” way, very environment-dependent
  - Not always used in modern implementations

# Query Phase

- Given a start  $q_s$  and goal  $q_g$ 
  1. Connect them to the roadmap using local planner
    - May need to try more than  $k$  nearest neighbors before connection is made
  2. Search  $G$  to find shortest path between  $q_s$  and  $q_g$  using A\*/Dijkstra's/etc.



# Path Shortening / Smoothing

- Don't even think of executing a path generated by a sampling-based planner without smoothing it!!!

## Shortcut Smoothing

For  $i = 0$  to MaxIterations


Pick two points,  $q_1$  and  $q_2$ , on the path randomly

Attempt to connect  $(q_1, q_2)$  with a line segment


If successful, replace path between  $q_1$  and  $q_2$  with the line segment




# Shortcut Smoothing



# Shortcut Smoothing




# Shortcut Smoothing




# PRM Failure Modes

1. Can't connect  $q_s$  and  $q_g$  to any nodes in the graph
  - Come up with an example in the graph below
2. Can't find a path in the graph but a path is possible
  - Come up with an example in the graph below



# Why do failures happen?

- Roadmap doesn't capture connectivity of space, to address this
  - Can run the learning phase longer
  - Can change sampling strategy to focus on narrow passages



- Local planner is too simple, to address this
  - Can use more sophisticated local planner

# Completeness

- Complete algorithms are slow.
  - A **complete** algorithm finds a path if one exists and reports no otherwise.
  - Example: Visibility graph
- Heuristic algorithms are unreliable.
  - Example: potential field
- **Probabilistic completeness**
  - Intuition: If there is a solution path, the algorithm will find it with high probability.

# Probabilistic Completeness

In an expansive space\*, the probability that a PRM planner fails to find a path when one exists goes to 0 exponentially in the number of milestones ( $\sim$  running time).

[Kavraki, Latombe, Motwani, Raghavan, 95]

[Hsu, Latombe, Motwani, 97]

\*Roughly, an expansive space is one where there are no infinitely-thin parts of free space.


# What happens in the limit for PRM?

- What if we ran the construction step of the PRM for infinite time...
  - What would the graph look like?
  - Would it capture the connectivity of the free space?
  - Would any collision-free start and goal be able to connect to the graph?
  - Is the PRM algorithm probabilistically complete?



Break


# PRM issues



- Two issues with the PRM:
  1. Uniform random sampling misses narrow passages
  2. Exploring whole space, but all we want is a path

# Sampling Strategies

- Most common is uniform random sampling
  - The bigger the area, the more likely it will be sampled
  - Problem: Narrow passages





- Are narrow passages inherently bad?
  - Does A\* running on a 2D grid have problems with narrow passages?

# OBPRM: An Obstacle-Based PRM

To Navigate Narrow Passages we must sample in them


- most PRM nodes are where planning is easy (not needed)



**Idea: Can we sample nodes near C-obstacle surfaces?**

- we cannot explicitly construct the C-obstacles...

## OBPRM: Finding Points on C-obstacles




### Basic Idea (for workspace obstacle S)

1. Find a point in S's C-obstacle  
(robot placement colliding with S)
2. Select a random direction in C-space
3. Find a free point in that direction
4. Find boundary point between them  
using binary search (collision checks)


Note: we can use more sophisticated heuristics to try to cover C-obstacle

# PRM vs OBPRM Roadmaps



## PRM

- 328 nodes
- 4 major CCs





## OBPRM

- 161 nodes
- 2 major CCs

# Sampling strategies: Gaussian


- Gaussian sampler
  - Pick a  $q_1$
  - Pick a  $q_2$  from a Gaussian distribution centered at  $q_1$
  - If **both** are in collision or collision-free, discard them, if one free, keep it




Sampling distribution for varying  $\sigma$   
(width decreasing from left to right)

# Sampling Strategies: Gaussian

- Performs well in narrow passages



Uniform Random Sampling



Gaussian Sampling




# Sampling Strategies


- Can we come up with a case where obstacle-biased sampling is worse than uniform random sampling?

# Sampling Strategies: Bridge

- Sample a  $q_1$  that is in collision
- Sample a  $q_2$  in neighborhood of  $q_1$  using some probability distribution (e.g. gaussian)
- If  $q_2$  in collision, get the midpoint of  $(q_1, q_2)$
- Check if midpoint is in collision, if not, add it as a node




# Sampling Strategies: Bridge



Gaussian  
Sampling

Bridge  
Sampling

What's going on  
at the corners?



Bridge Sampling performs  
well in narrow passages

# Rapidly-exploring Random Trees (RRTs)

---

# Single-query methods

- Motivation: Why try to capture the connectivity of the whole space when all you need is one path?
- Algorithms:
  - Single-Query BiDirectional Lazy PRM (SBL-PRM)
  - Expansive Space Trees (EST)
  - **Rapidly-exploring Random Tree (RRT)**
    - AKA “RDT” in the book
- Key idea: Build a *tree* instead of a general graph.
- The tree grows in  $C_{free}$ 
  - Like PRM, captures some connectivity
  - Unlike PRM, only explores what is connected to  $q_{start}$

# Naïve Tree Algorithm


$q_{\text{node}} = q_{\text{start}}$

For  $i = 1$  to NumberSamples


$q_{\text{rand}} = \text{Sample near } q_{\text{node}}$

    Add edge  $e = (q_{\text{rand}}, q)$  if  
collision-free


$q_{\text{node}} = \text{Pick random node of tree}$



# RRT Growing in Empty Space



RRT with obstacles and goal bias.






# Path Planning with Rapidly-Exploring Random Trees (RRTs)

```
BUILD_RRT( $q_{init}$ ) {  
   $T.init(q_{init})$ ;  
  for  $k = 1$  to  $K$  do  
     $q_{rand} = \text{RANDOM\_CONFIG}()$ ;  
     $\text{EXTEND}(T, q_{rand})$   
}
```

$\text{EXTEND}(T, q_{rand})$




# RRT Goal Biasing

- In “pure” form RRTs are great at filling space, but we need a path!
- Need to bias RRTs toward goal to produce a path
  - When generating a random sample, with some probability pick the goal instead of a random node
  - This introduces another parameter
  - James Kuffner’s experience is that 5-10% is the right choice
- What happens if you set probability of sampling goal to 100%?

# RRT Extension Types

- RRT-Extend
  - Take one step toward a random sample
- RRT-Connect
  - Step toward random sample until it is either
    - Reached
    - You hit an obstacle




# BiDirectional RRTs

- BiDirectional RRT
  - Grow trees from both start and goal
  - Try to get trees to connect to each other
  - Trees can both use Extend or both use Connect or one use Extend and one Connect
- BiDirectional RRT with Connect for both trees is my favorite, I always try this first
  - This variant has only one parameter; the step size

# Example of BiDirectional RRT



Connect




Extend

# 1) One tree grown using random target




Connect




Extend

## 2) New node becomes target for other tree




Connect



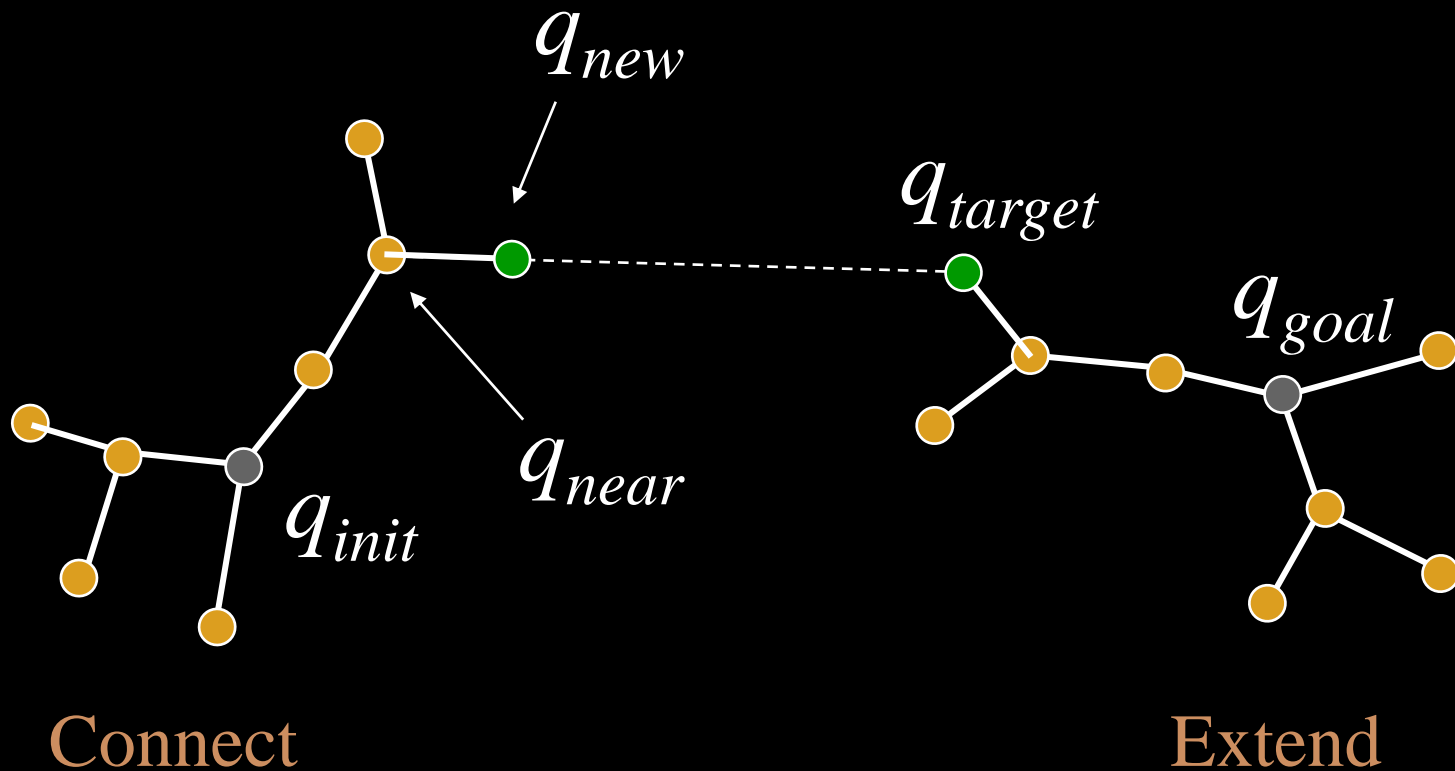
Extend

### 3) Calculate node “nearest” to target







## 4) Try to add new collision-free branch




5) If successful, keep extending branch




5) If successful, keep extending branch




5) If successful, keep extending branch



## 6) Path found if branch reaches target




## 7) Return path connecting start and goal




# Path Smoothing/Optimization

- RRTs produce notoriously bad paths
  - Not surprising since no consideration of path quality
- ALWAYS smooth/optimize the returned path
  - Many methods exists, e.g. shortcut smoothing



# RRT Examples: The Alpha Puzzle



- VERY hard 6DOF motion planning problem (long, winding narrow passage)



- *“In 2001, it was solved by using a balanced bidirectional RRT, developed by James Kuffner and Steve LaValle. There are no special heuristics or parameters that were tuned specifically for this problem. On a current PC (circa 2003), it consistently takes a few minutes to solve” –RRT website*
- RRT became famous in large part because it was able to solve this puzzle



# RRT Examples: Articulated Objects



# RRT Analysis

The limiting distribution of vertices:

- **THEOREM:**  $\mathbf{X}_k$  converges to  $\mathbf{X}$  with probability 1 as time goes to infinity

$\mathbf{X}_k$  : The RRT vertex distribution at iteration  $k$

$\mathbf{X}$  : The distribution used for generating samples

- If using uniform distribution, tree nodes converge to the free space
- Based on this, we can prove that RRT is probabilistically complete

# Summary: Sampling-Based Planning

- The good:
  - Provides fast *feasible* solution
  - Popular methods have few parameters
  - Works on practical problems
  - Works in high-dimensions
  - Works even with the wrong distance metric

# Summary: Sampling-Based Planning

- The bad:
  - No quality guarantees on paths\*
    - In practice: smooth/optimize path afterwards
  - No termination when there is no solution
    - In practice: set an arbitrary timeout
  - Probabilistic completeness is a weak property
    - Completeness in high-dimensions is impractical

\*More recent methods have asymptotic optimality guarantees (e.g. RRT\*)

# Homework

- LaValle Ch. 14-14.5