

Search for Optimization



Real-World Problems

- *“many real world problems have a landscape that looks more like a widely scattered family of balding porcupines on a flat floor, with miniature porcupines living on the tip of each porcupine needle, ad infinitum.”*

-Russell and Norvig

A different view of optimization

- In lectures on convex optimization, we assumed the function was convex, now we remove that assumption!



- Removing this allows us to consider
 - non-convex continuous problems
 - problems where variables are discrete (really important!)
- Unfortunately, not much we can prove in non-convex optimization ☹️

Outline


- Graphs
- Local search methods
 - Hill Climbing
- Simulated Annealing
- Genetic Algorithms
- Genetic Programming

Graphs

- A **graph** is a set of **vertices** (also called **nodes**) V and **edges** E :

$$G = (V, E)$$

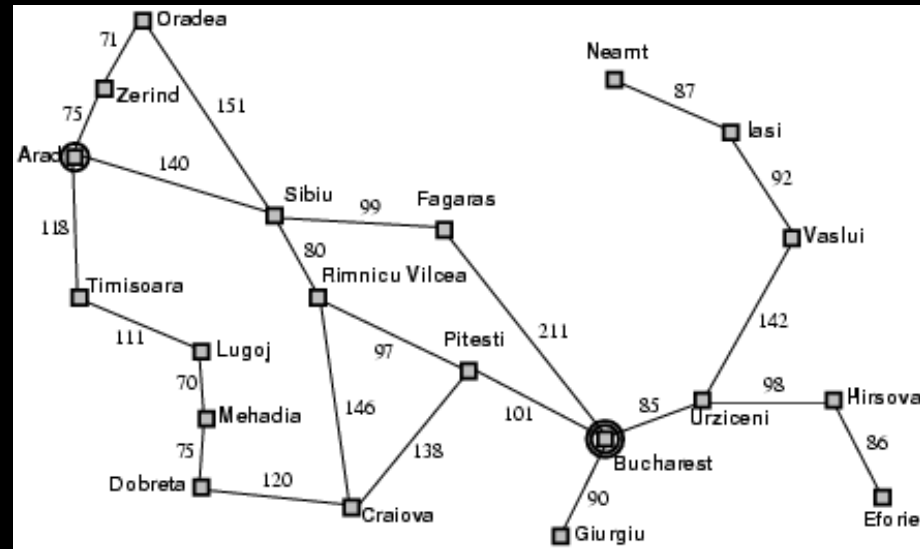
- Example:



- $V = \{1, 2, 3, 4, 5, 6\}$
- $E = \{(6, 4), (4, 5), (4, 3), (3, 2), (5, 2), (2, 1), (5, 1)\}$
- In general there may be an infinite number of vertices and edges


Graphs: Why this representation?

- Graphs capture the idea of **adjacency** – i.e. what is “next to” what
 - A node u 's **neighbors** are the set of nodes connected to u by an edge
- We can use adjacency relationships to search the graph for a certain node or a path between nodes
- Example:



Graphs for optimization

- For optimization, adjacency between nodes can be used to determine what solutions to explore next
- Consider each node in the graph to represent a solution to a problem:




- Here solution 1 is adjacent to solution 2 because they are “close” to each other in solution space
- Adjacency must be defined by the algorithm designer
 - A good adjacency definition helps you search the space of solutions systematically


Example: n -queens

- The problem: Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

An example solution (not good!):




A much better solution:



- Define adjacency: Two solutions u and v are adjacent iff $n - 1$ queens are at the same position in u and v
 - You could make up other definitions, too!

Example: n -queens

- Starting with a bad solution, can use adjacency to search for a good one:



Local search algorithms

- To try to solve optimization problems using *search* we
 1. Define the set of possible solutions
 - This is then the set of nodes
 2. Define adjacency between solutions
 - This defines the set of edges
 3. Define a cost/value/fitness function for nodes
 - Outputs how good a solution is
- Can use **local search algorithms** to search the graph for the best solution
 - keep a (sometimes) single "current" state, try to improve it
- Descent methods (from before) are a form of local search algorithms
 - But these only work for convex continuous functions
- Some local search methods can also handle discrete variables

Local Search

- Operates by keeping track of only the current solution and moving only to neighbors of that solution
- Often used for:
 - Optimization problems
 - Scheduling
 - Task assignment
 - ...many other problem where the goal is to find the best state according to some **objective function**

Local Search: Hill-climbing

Hill-climbing search

- Consider next possible moves (i.e. neighbors)
 - Pick the one that improves cost/value/fitness function the most
- “Like climbing Everest in thick fog with amnesia”

Hill-climbing search

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

inputs: *problem*, a problem

local variables: *current*, a node

neighbor, a node

current ← MAKE-NODE(INITIAL-STATE[*problem*])

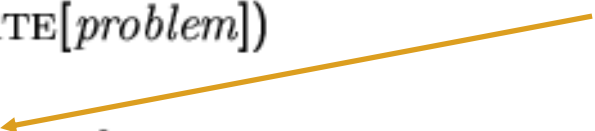
loop do

neighbor ← a highest-valued successor of *current*









if VALUE[*neighbor*] ≤ VALUE[*current*] **then return** STATE[*current*]

current ← *neighbor*

“successor” is
a synonym for
neighbor




Hill-climbing search: 8-queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14		13	16	13	16
	14	17	15		14	16	16
17		16	18	15		15	
18	14		15	15	14		16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state

Hill-climbing search: 8-queens problem


- 5 steps later...



- A **local minimum** with $h = 1$ (a common problem with hill climbing)

Drawbacks of hill climbing

- Problem: depending on initial state, can get stuck in local maxima
 - I.e. where you end depends on where you start




Hill Climbing with Local Minima: Try, try again

- Run algorithm some number of times and return the best solution
 - Initial start location is usually chosen randomly
- If you run it “enough” times, will get answer (in the limit)
- Drawback: takes lots of time, never sure when to terminate

Simulated Annealing

Simulated annealing

- Hill climbing problems
 - Gets stuck on plateaus
 - Returns sub-optimal solutions (local minima)
- Simulated annealing main idea: explicitly inject variability into the search process



Properties of simulated annealing

- More variability at the beginning of search
 - Since you have little confidence you're in right place
- Variability decreases over time
 - Don't want to move away from a good solution
- Probability of picking a move is related to how good it is
 - No decrease or slight decrease are more likely than major decreases in quality

Simulated annealing implementation

- Create a “temperature schedule” for how variability changes as we iterate

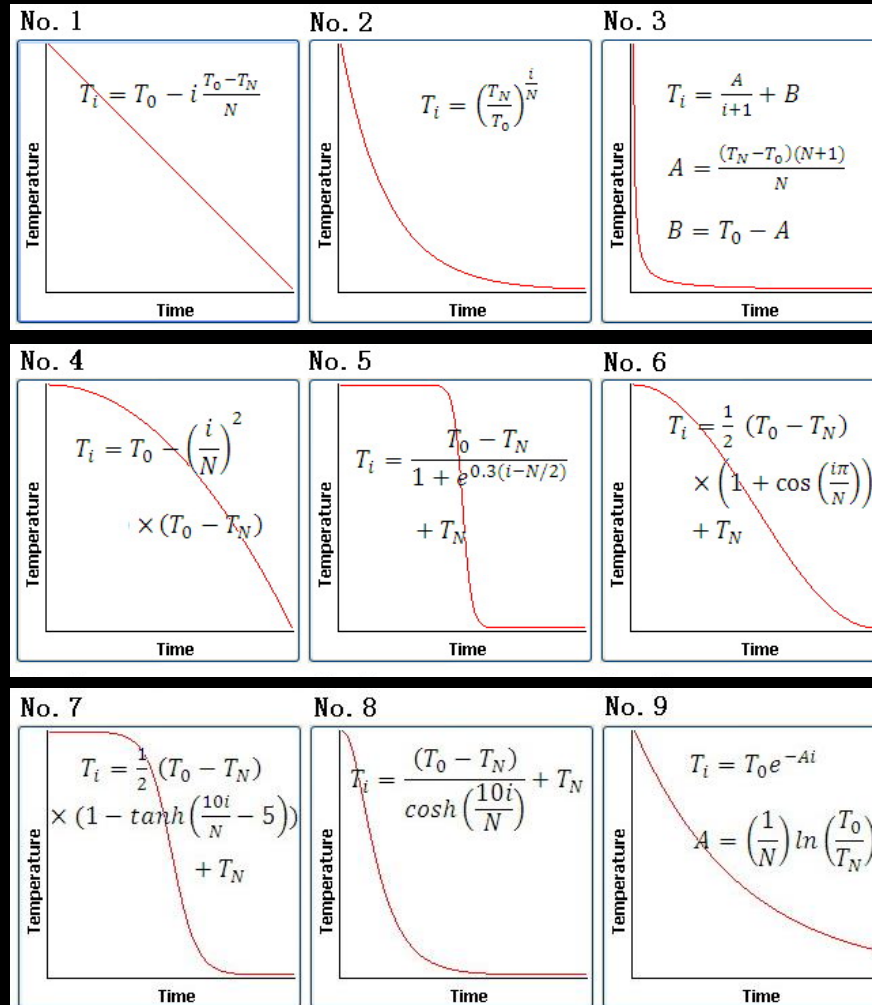
```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Much faster than
considering all neighbors
(in high dimensions)

Temperature schedules


- Which one to use?



Evolutionary Algorithms

Genetic algorithms

- Inspired (loosely) by the process of evolution in nature
- Operators:
 - **Crossover**: New states generated from two *parent* states
 - **Mutation**: Randomly change a component of a state




Genetic Algorithms

1. Initialize **population** (k random states)
2. Select a set of parents from the population for mating (based on fitness)
3. Generate children via crossover of parents
4. Mutation (add randomness to the children's variables)
5. Evaluate fitness of children
6. Replace worst parents with the children
7. Go to step 2


Genetic algorithms

24748552
32752411
24415124
32543213
(a)
Initial Population



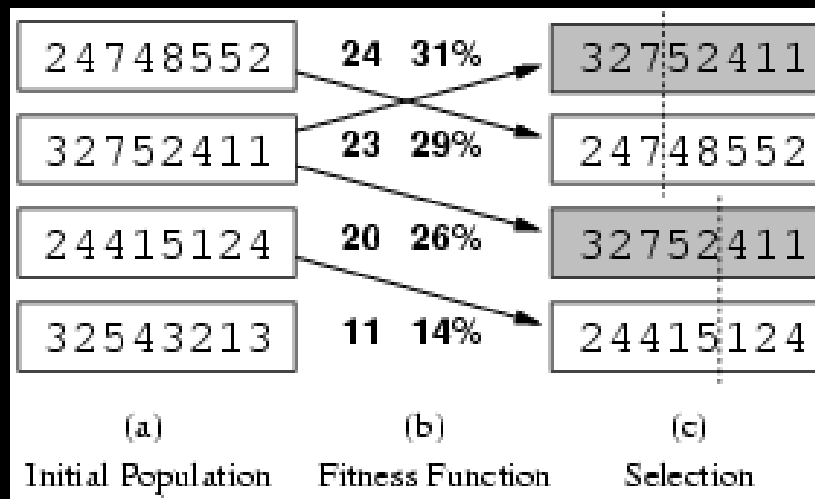
32752411

Genetic algorithms



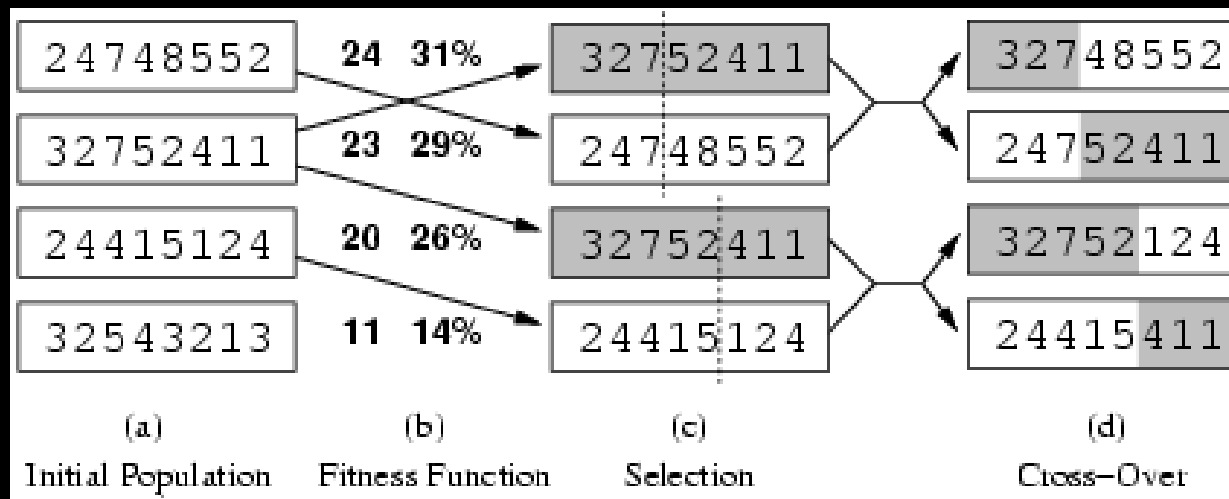
- Fitness function: number of *non-attacking* pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
- Compute probability of being selected from population:
 - $24/(24+23+20+11) = 31\%$
 - $23/(24+23+20+11) = 29\%$
 - ... etc.

Genetic algorithms




Select individuals based on probabilities

Genetic algorithms

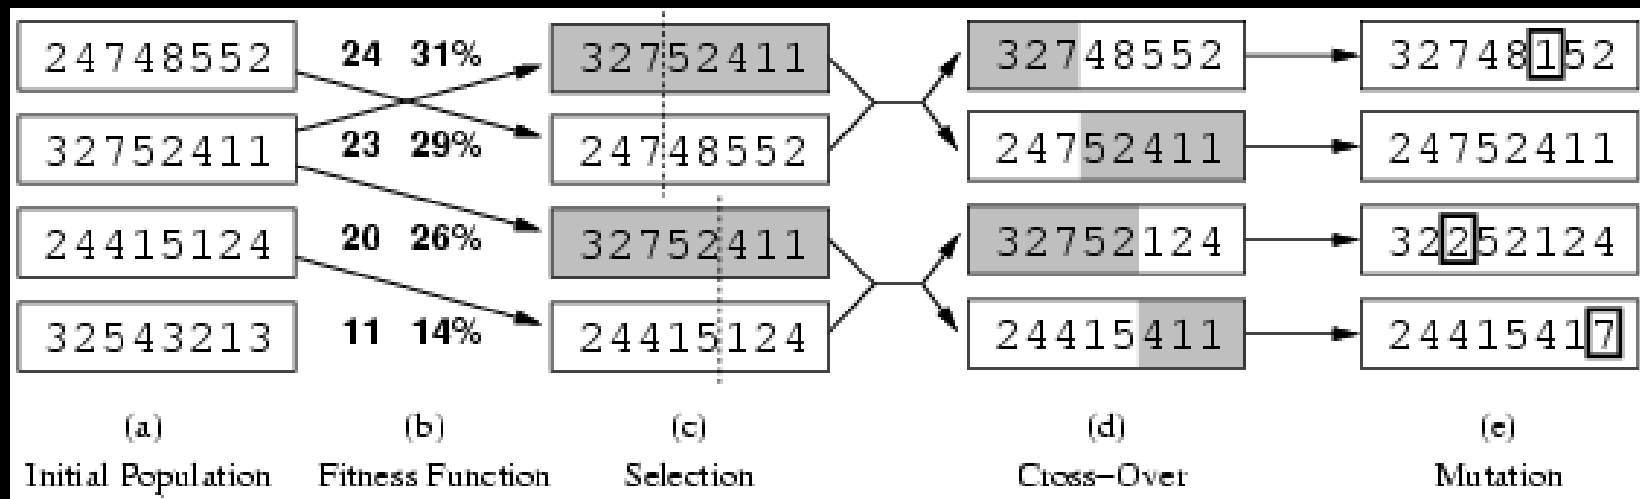


Select individuals based on probabilities



Crossover from the top two parents.


Genetic algorithms



Break

Genetic Programming

- Evolve *functions* instead of vectors of numbers
- Individuals represented as trees:




Non-leaf nodes are operators

Leaf nodes are operands
(input variables or constants)


- Useful for regression problems

Genetic Programming

- Crossover: Swap branches of the tree




- Mutation: Replace a node with a random node or replace an input to a node




- The hard part: how do you pick your set of operators?
- Matlab toolbox: GPLAB
- Python: DEAP

Application of Genetic Programming: Automating Science



Application of Genetic Programming: Automating Science



The Power of Evolutionary Algorithms

- Human-competitive genetic programming:
 - <http://www.genetic-programming.com/humancompetitive.html>
- Evolving robot morphology and walking:
 - <http://www.uvm.edu/~uvmpr/?Page=news&storyID=11482&category=ucommfeatureb>
- Evolving soft robots:
 - <https://www.youtube.com/watch?v=z9ptOeByLA4>

Summary

- Can represent optimization problems as graphs where solutions are nodes and adjacency of solutions determines edges
- Local search (like hill-climbing) often finds local minima
- Can inject randomness to avoid getting stuck in local minima
 - Simulated Annealing
 - Genetic Algorithms
 - Genetic Programming

Homework

- HW 2 is posted (START EARLY)
- AI book Ch. 3.1-3.5.2
- LaValle Ch. 2.0 - 2.3