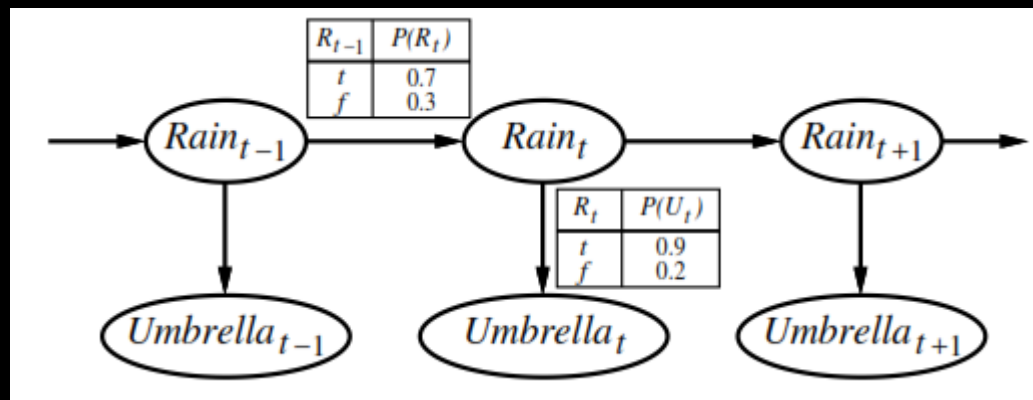


Bayes Filter and Kalman Filters

Using materials from probabilistic-robotics.org, AIMA book

Last time...

- We saw how to incorporate time into probabilistic reasoning (in the form of a Bayes net)
- We made the **Markov Assumption** to keep the inference manageable



- But we only considered cases where the state was discrete
- Today we will look at algorithms that handle continuous distributions

Outline

- Hidden Markov Models (HMMs) (review)
- Bayes filter
- Kalman Filter
- Extended Kalman Filter (EKF)

Inference Tasks

Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Prediction: $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$ for $k > 0$

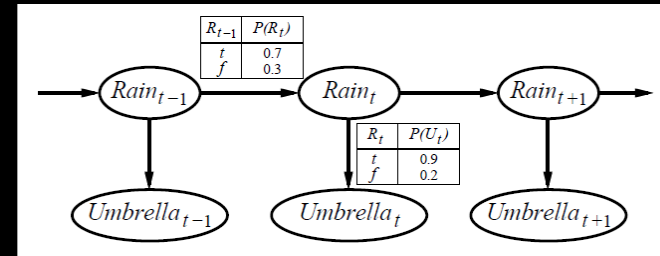
evaluation of possible action sequences;
like filtering without the evidence

Smoothing: $P(\mathbf{X}_k | \mathbf{e}_{1:t})$ for $0 \leq k < t$

better estimate of past states, essential for learning

Most likely explanation: $\arg \max_{\mathbf{x}_{1:t}} P(\mathbf{x}_{1:t} | \mathbf{e}_{1:t})$

speech recognition, decoding with a noisy channel



Hidden Markov Models (HMMs) (review)

- An HMM is a temporal probabilistic model in which the state is described by a **single discreet random variable**.

X_t is a single, discrete variable (usually E_t is too)

Domain of X_t is $\{1, \dots, S\}$

Transition matrix $T_{ij} = P(X_t = j | X_{t-1} = i)$, e.g., $\begin{pmatrix} 0.7 & 0.3 \\ 0.3 & 0.7 \end{pmatrix}$

Sensor matrix O_t for each time step, diagonal elements $P(e_t | X_t = i)$

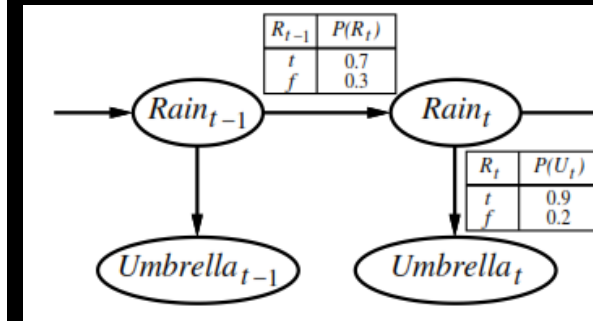
e.g., with $U_1 = \text{true}$, $O_1 = \begin{pmatrix} 0.9 & 0 \\ 0 & 0.2 \end{pmatrix}$

Forward and backward messages as column vectors:

$$\mathbf{f}_{1:t+1} = \alpha \mathbf{O}_{t+1} \mathbf{T}^\top \mathbf{f}_{1:t}$$

$$\mathbf{b}_{k+1:t} = \mathbf{T} \mathbf{O}_{k+1} \mathbf{b}_{k+2:t}$$


Forward-backward algorithm needs time $O(S^2t)$ and space $O(St)$



Bayes Filter

HMM vs. Bayes Filter


- In HMMs, the system is passive: just a stream of perception data
- A robot can take *actions* u as well as get perceptions z



- Robot state and perception data are usually multi-dimensional
- When the state is discrete, we can use a **Bayes filter**

Bayes Filter Belief

- The state must be discrete, so we usually use a grid to represent it
- Each grid cell contains the belief $\text{Bel}(x_t)$ (the probability that the true state of the system is x_t)
- E.g. for a mobile robot:



Discrete Bayes Filter Algorithm

- Given a piece of sensor or action data, update $Bel(x)$ using this algorithm:




```
1.  Algorithm Discrete_Bayes_filter(  $Bel(x), d$  ):  
2.   $\eta = 0$   
3.  If  $d$  is a perceptual data item  $z$  then  
4.  For all  $x$  do  
5.       $Bel'(x) = P(z | x) Bel(x)$   
6.       $\eta = \eta + Bel'(x)$   
7.  For all  $x$  do  
8.       $Bel'(x) = \underline{\eta}^{-1} Bel'(x)$   
9.  Else if  $d$  is an action data item  $u$  then  
10. For all  $x$  do  
11.      $Bel'(x) = \sum_{x'} P(x | u, x') \underline{Bel(x')}$   
12. Return  $Bel'(x)$ 
```

Very inefficient
for large state
spaces!
There are
heuristics
which only
update the
belief locally




Grid-based Localization Example

Perception data




Belief



Kalman Filter

Kalman Filters

- The real world is not discrete! Need to consider **continuous** variables



- Kalman filters are used to track state of robots, chemical plants, planets, etc.
- **Key Idea:** Arbitrary continuous models are intractable, so represent everything with *Gaussians*
 - Gaussian prior, linear Gaussian transition model and sensor model


Gaussians

- Univariate

$$p(x) \sim N(\mu, \sigma^2):$$

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}}$$


mean *variance*



- Multivariate

$$p(\mathbf{x}) \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}):$$

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2} (\mathbf{x}-\boldsymbol{\mu})^t \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu})}$$



Multivariate Gaussians

$$\left. \begin{array}{l} X \sim N(\mu, \Sigma) \\ Y = AX + B \end{array} \right\} \Rightarrow Y \sim N(A\mu + B, A\Sigma A^T)$$

matrix *vector*

$$\left. \begin{array}{l} X_1 \sim N(\mu_1, \Sigma_1) \\ X_2 \sim N(\mu_2, \Sigma_2) \end{array} \right\} \Rightarrow p(X_1) \cdot p(X_2) \sim N\left(\frac{\Sigma_2}{\Sigma_1 + \Sigma_2} \mu_1 + \frac{\Sigma_1}{\Sigma_1 + \Sigma_2} \mu_2, \frac{1}{\Sigma_1^{-1} + \Sigma_2^{-1}}\right)$$

We stay in the “Gaussian world” as long as we start with Gaussians and perform only linear transformations.

Kalman Filter

Estimates the state x of a discrete-time controlled process that is governed by the linear stochastic difference equation

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

Current
state

Previous
state

Action

Actuation
noise

with a sensor measurement

$$z_t = C_t x_t + \delta_t$$

Current
measurement

Current
state

Sensor
noise

Example for GPS sensor:

$$z_t = \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} pos_x \\ vel_x \end{bmatrix} + \delta_t$$

Components of a Kalman Filter

A_t

Matrix ($n \times n$) that describes how the state changes from $t-1$ to t without controls or noise.

B_t

Matrix ($n \times |u|$) that describes how the control u_t changes the state from $t-1$ to t .

C_t

Matrix ($k \times n$) that describes how to map the state x_t to an observation z_t .

ε_t

Random variables representing the process and measurement noise that are assumed to be independent and normally distributed

δ_t


with covariance R_t and Q_t respectively.

Kalman Filter State Estimate

- The Kalman filter computes a Gaussian probability distribution of the state given the prior distribution and a sensor measurement

$$p(\mathbf{x}) \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) :$$


$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^t \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$



- A Gaussian is represented by $\boldsymbol{\mu}$ (a vector) and $\boldsymbol{\Sigma}$ (a matrix), so all we need to do is track these two variables



Kalman Filter Algorithm

Algorithm **Kalman_filter**(μ_{t-1} , Σ_{t-1} , u_t , z_t):




See here for derivation. Note that they use different notation!


The Prediction-Correction-Cycle




$$\begin{aligned}\bar{\mu}_t &= A_t \mu_{t-1} + B_t u_t \\ \bar{\Sigma}_t &= A_t \Sigma_{t-1} A_t^T + \textcircled{R_t}\end{aligned}$$



The Prediction-Correction-Cycle




$$\mu_t = \bar{\mu}_t + K_t(z_t - C_t\bar{\mu}_t)$$
$$\Sigma_t = (I - K_tC_t)\bar{\Sigma}_t$$



Correction

Prediction-Correction Cycle



Kalman Filter Localization Example



Blue line: actual trajectory
Blue dots: GPS measurements
Red line: estimated states

by Keyan Ghazi-Zahedi
https://www.youtube.com/watch?v=ZYexI6_zUMkby

Kalman Filter Highlights

- **Highly efficient**: Polynomial in measurement dimensionality k and state dimensionality n :

$$O(k^{2.376} + n^2)$$

- **Optimal for linear Gaussian systems!**
(known R and Q)
- But, most robotics systems are **nonlinear**! ☹️

BREAK

Extended Kalman Filter (EKF)

Nonlinear dynamic systems

- Most robotics problems involve nonlinear dynamics and sensors

$$x_t = g(u_t, x_{t-1})$$

$$z_t = h(x_t)$$

The EKF trick

- Can't deal with non-linear functions directly
- But, if the change is small, we can use a *local linear approximation*
- How? Compute the Jacobians of g and h !

$$x_t = \underline{g(u_t, x_{t-1})} \approx g(u_t, \underline{\mu_{t-1}}) + \frac{\partial g(u_t, \mu_{t-1})}{\partial x_{t-1}} (\underline{x_{t-1}} - \underline{\mu_{t-1}})$$
$$x_t = g(u_t, x_{t-1}) \approx g(u_t, \mu_{t-1}) + G_t (x_{t-1} - \mu_{t-1})$$

$$G_t = \frac{\partial g(u_t, \mu_{t-1})}{\partial x_{t-1}}$$

$$z_t = h(\underline{x_t}) \approx h(\underline{\bar{\mu}_t}) + \frac{\partial h(\bar{\mu}_t)}{\partial x_t} (x_t - \bar{\mu}_t)$$
$$z_t = h(x_t) \approx h(\bar{\mu}_t) + H_t (x_t - \bar{\mu}_t)$$

$$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t}$$

Algorithm **Extended_Kalman_filter**($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$):

1. Prediction:

$$\bar{\mu}_t = g(u_t, \mu_{t-1})$$

$$\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$$

2. Correction:

$$K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t))$$

$$\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$$

Kalman Filter

$$\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$$

$$K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$


$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

Return μ_t, Σ_t

$$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t}$$

$$G_t = \frac{\partial g(u_t, \mu_{t-1})}{\partial x_{t-1}}$$

Example: Beacon-based Robot Localization



Example Motion Model

- State is $x_t = (x_t, y_t, \theta_t)$
- Command is rotation, translation, rotation

$$u_t = (\delta_{rot_1}, \delta_{trans}, \delta_{rot_2})$$

- Actual motion is $(\tilde{\delta}_{rot_1}, \tilde{\delta}_{trans}, \tilde{\delta}_{rot_2})$, a noisy version of the command


- Motion model g is:

$$x_{t+1} = x_t + \tilde{\delta}_{trans} \cos(\theta_t + \tilde{\delta}_{rot_1})$$

$$y_{t+1} = y_t + \tilde{\delta}_{trans} \sin(\theta_t + \tilde{\delta}_{rot_1})$$

$$\theta_{t+1} = \theta_t + \tilde{\delta}_{rot_1} + \tilde{\delta}_{rot_2}$$

(modulo 2π)




Not linear!

Example sensor model

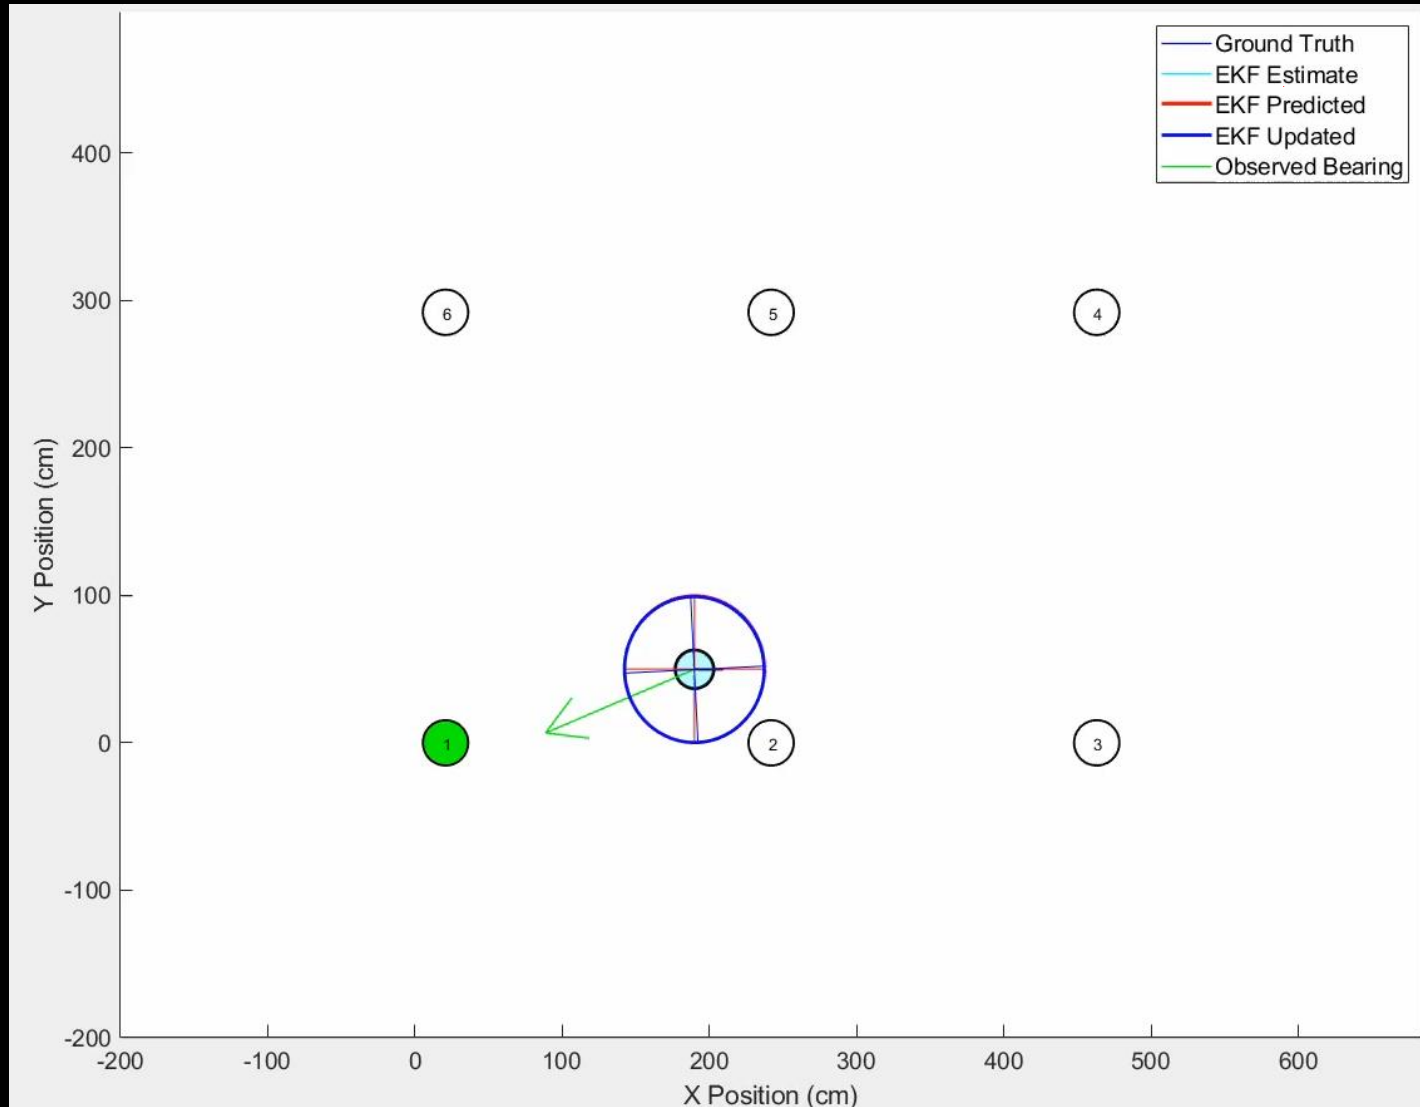
- The map is known
 - Beacons are at known positions
- Sensor reports noisy bearing $\tilde{\theta}$ and exact landmark ID L
 - Only one beacon is observed at one time

$$z_t = \begin{pmatrix} \tilde{\theta} \\ L \end{pmatrix} = \begin{pmatrix} \text{atan2}(y_{rob} - y_L, x_{rob} - x_L) \\ L \end{pmatrix}$$

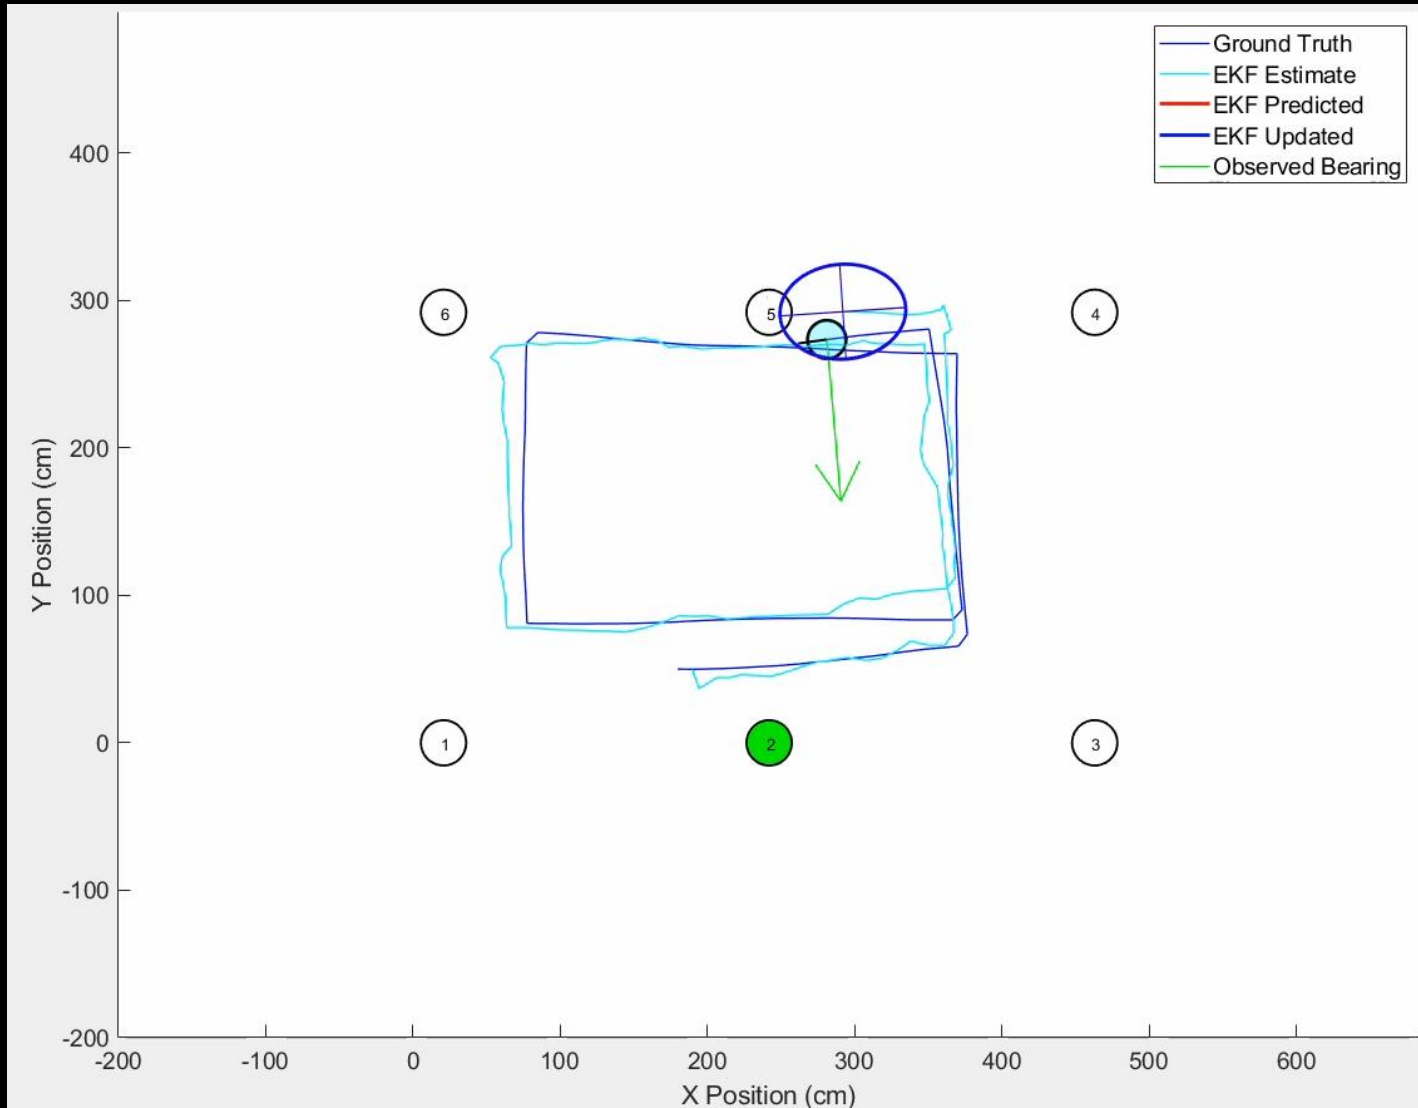
Not linear!




EKF Localization Example



What happened here?



EKF Inaccuracy



- Robot is close to beacon, so small changes in the position cause large changes in the measurement
- Observation Jacobian H_t has large terms
 - Causes filter to become confident in position
- But, measurement error is large: -1.86 radians!
 - Causes filter to be wrong about position due to non-linearities

EKF Highlights


- **Highly efficient**: Polynomial in measurement dimensionality k and state dimensionality n :

$$O(k^{2.376} + n^2)$$

- **Not optimal!**
- Can **diverge** if nonlinearities are large!
- In practice, works surprisingly well even when assumptions are violated

Overall Kalman Filter Limitations

- Greatest strength is the greatest weakness
 - Everything must be a Gaussian!
- Cannot be used if transition is non-linear:



Summary

- HMMs are Bayes nets that make the Markov assumption for a single **discrete** random variable
- Bayes Filter tracks a discrete state
 - Computationally intractable in high dimensions
- Kalman Filters track **continuous** random state variables for linear systems
 - Can have multiple state and sensor variables
 - Assume state and measurement distributions are Gaussian
 - Fast updates at the expense of generality
- Extended Kalman Filter (EKF) is for nonlinear systems
 - Still using Gaussian state and measurement distributions
 - Use Jacobian of motion model and observation model functions
 - Works well when function is close to linear locally

Homework

- Particle Filters, Sec. 2.1