


# Probability with time and HMMs

---

Some examples from Sebastian Thrun

## Previously...

- We talked about Bayes nets as ways to represent probability distributions



- But what if the events we want to represent change over time?

# Outline

- Time and uncertainty
- Inference: filtering, prediction, smoothing
- Hidden Markov Models

## Time and uncertainty

The world changes; we need to track and predict it

Basic idea: copy state and evidence variables for each time step

$\mathbf{X}_t$  = set of unobservable state variables at time  $t$   
e.g., *BloodSugar<sub>t</sub>*, *StomachContents<sub>t</sub>*, etc.

$\mathbf{E}_t$  = set of observable evidence variables at time  $t$   
e.g., *MeasuredBloodSugar<sub>t</sub>*, *PulseRate<sub>t</sub>*, *FoodEaten<sub>t</sub>*

This assumes **discrete time**; step size depends on problem

Notation:  $\mathbf{X}_{a:b} = \mathbf{X}_a, \mathbf{X}_{a+1}, \dots, \mathbf{X}_{b-1}, \mathbf{X}_b$


# Markov processes (Markov chains)

Construct a Bayes net from these variables: parents?

Markov assumption:  $\mathbf{X}_t$  depends on **bounded** subset of  $\mathbf{X}_{0:t-1}$

First-order Markov process:  $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$

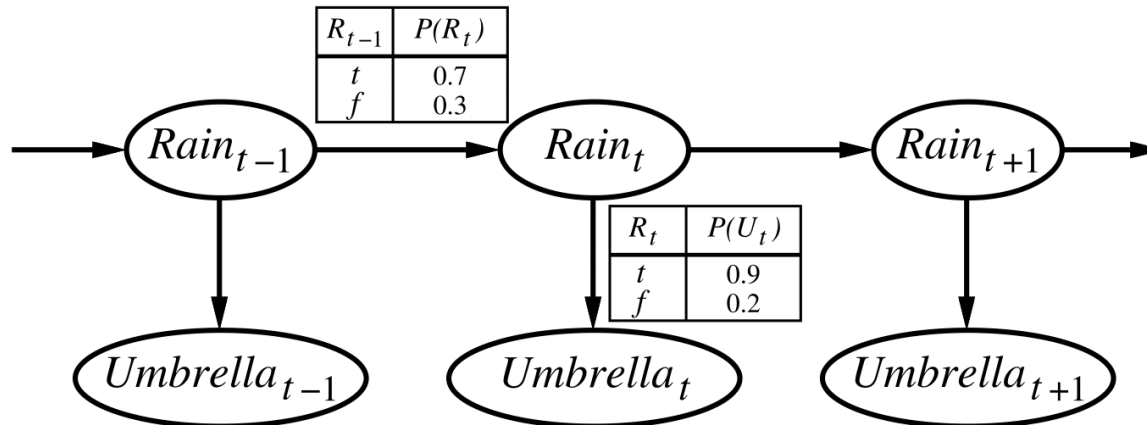
Second-order Markov process:  $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-2}, \mathbf{X}_{t-1})$



Sensor Markov assumption:  $\mathbf{P}(\mathbf{E}_t | \mathbf{X}_{0:t}, \mathbf{E}_{0:t-1}) = \mathbf{P}(\mathbf{E}_t | \mathbf{X}_t)$

Stationary process: transition model  $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$  and sensor model  $\mathbf{P}(\mathbf{E}_t | \mathbf{X}_t)$  fixed for all  $t$

## Example



First-order Markov assumption not exactly true in real world!

Possible fixes:

1. **Increase order** of Markov process
2. **Augment state**, e.g., add  $Temp_t$ ,  $Pressure_t$

Example: robot motion.

Augment position and velocity with  $Battery_t$

## Inference tasks

Filtering:  $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Prediction:  $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$  for  $k > 0$

evaluation of possible action sequences;  
like filtering without the evidence

Smoothing:  $P(\mathbf{X}_k | \mathbf{e}_{1:t})$  for  $0 \leq k < t$

better estimate of past states

Most likely explanation:  $\arg \max_{\mathbf{x}_{1:t}} P(\mathbf{x}_{1:t} | \mathbf{e}_{1:t})$

speech recognition, decoding with a noisy channel

# Filtering

Aim: devise a **recursive** state estimation algorithm:

$$\mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t+1}) = f(\mathbf{e}_{t+1}, \mathbf{P}(\mathbf{X}_t|\mathbf{e}_{1:t}))$$

$$\begin{aligned}\mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t+1}) &= \mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t}, \mathbf{e}_{t+1}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1}|\mathbf{X}_{t+1}, \mathbf{e}_{1:t}) \mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1}|\mathbf{X}_{t+1}) \mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t})\end{aligned}$$

I.e., **prediction** + **estimation**. Prediction by summing out  $\mathbf{X}_t$ :

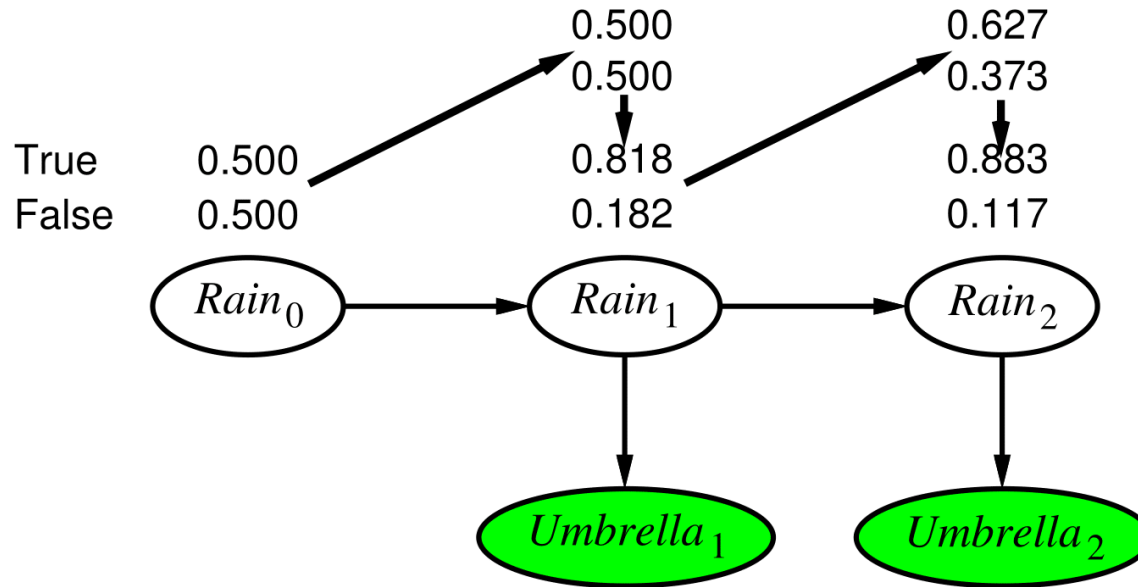
$$\begin{aligned}\mathbf{P}(\mathbf{X}_{t+1}|\mathbf{e}_{1:t+1}) &= \alpha \mathbf{P}(\mathbf{e}_{t+1}|\mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1}|\mathbf{x}_t, \mathbf{e}_{1:t}) P(\mathbf{x}_t|\mathbf{e}_{1:t}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1}|\mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1}|\mathbf{x}_t) P(\mathbf{x}_t|\mathbf{e}_{1:t})\end{aligned}$$

$\mathbf{f}_{1:t+1} = \text{FORWARD}(\mathbf{f}_{1:t}, \mathbf{e}_{t+1})$  where  $\mathbf{f}_{1:t} = \mathbf{P}(\mathbf{X}_t|\mathbf{e}_{1:t})$


Time and space **constant** (independent of  $t$ )



## Filtering example



# Smoothing



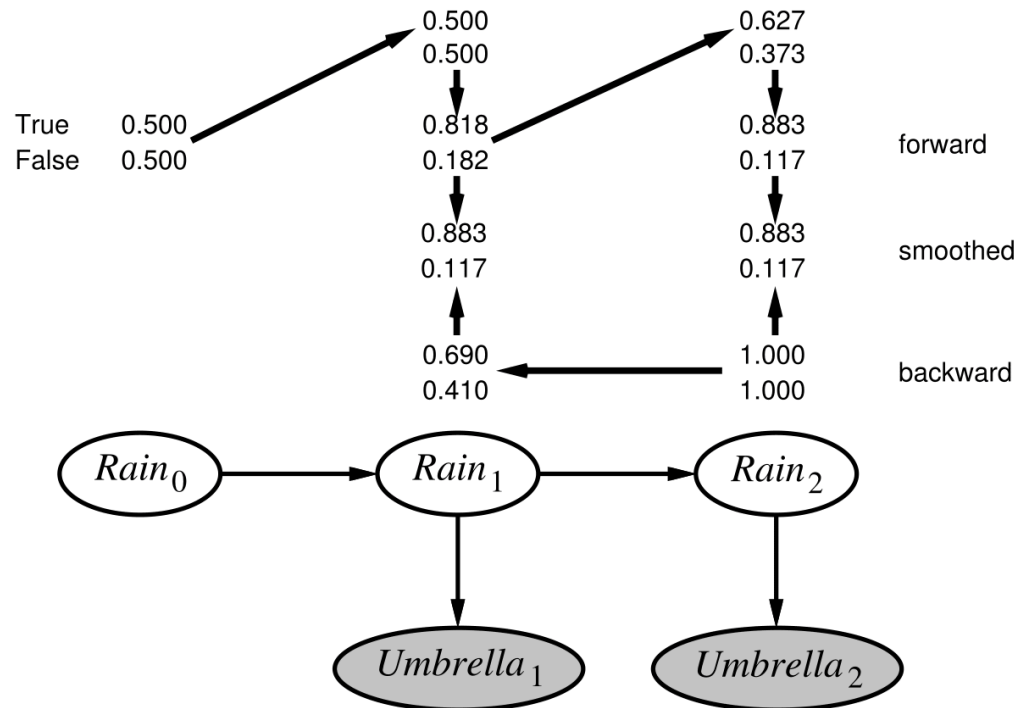
Divide evidence  $\mathbf{e}_{1:t}$  into  $\mathbf{e}_{1:k}$ ,  $\mathbf{e}_{k+1:t}$ :

$$\begin{aligned}
 \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:t}) &= \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\
 &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{e}_{1:k}) \\
 &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) \\
 &= \alpha \mathbf{f}_{1:k} \mathbf{b}_{k+1:t}
 \end{aligned}$$

Backward message computed by a backwards recursion:

$$\begin{aligned}
 \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k)
 \end{aligned}$$

# Smoothing example



Forward-backward algorithm: cache forward messages along the way  
 Time linear in  $t$  (polytree inference), space  $O(t|f|)$

## Most likely explanation

Most likely sequence  $\neq$  sequence of most likely states!!!!

Most likely path to each  $\mathbf{x}_{t+1}$

= most likely path to **some**  $\mathbf{x}_t$  plus one more step

$$\begin{aligned} & \max_{\mathbf{x}_1 \dots \mathbf{x}_t} \mathbf{P}(\mathbf{x}_1, \dots, \mathbf{x}_t, \mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \\ &= \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \max_{\mathbf{x}_t} \left( \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t) \max_{\mathbf{x}_1 \dots \mathbf{x}_{t-1}} P(\mathbf{x}_1, \dots, \mathbf{x}_{t-1}, \mathbf{x}_t | \mathbf{e}_{1:t}) \right) \end{aligned}$$

Identical to filtering, except  $\mathbf{f}_{1:t}$  replaced by


$$\mathbf{m}_{1:t} = \max_{\mathbf{x}_1 \dots \mathbf{x}_{t-1}} \mathbf{P}(\mathbf{x}_1, \dots, \mathbf{x}_{t-1}, \mathbf{X}_t | \mathbf{e}_{1:t}),$$

I.e.,  $\mathbf{m}_{1:t}(i)$  gives the probability of the most likely path to state  $i$ .

Update has sum replaced by max, giving the **Viterbi algorithm**:

$$\mathbf{m}_{1:t+1} = \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \max_{\mathbf{x}_t} (\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t) \mathbf{m}_{1:t})$$

# Viterbi example



# Hidden Markov Models HMMs

---

# Hidden Markov Models (HMMs)

- An HMM is a temporal probabilistic model in which the state is described by a **single discrete random variable**.

$X_t$  is a single, discrete variable (usually  $E_t$  is too)

Domain of  $X_t$  is  $\{1, \dots, S\}$

Transition matrix  $T_{ij} = P(X_t = j | X_{t-1} = i)$ , e.g.,  $\begin{pmatrix} 0.7 & 0.3 \\ 0.3 & 0.7 \end{pmatrix}$

Sensor matrix  $O_t$  for each time step, diagonal elements  $P(e_t | X_t = i)$

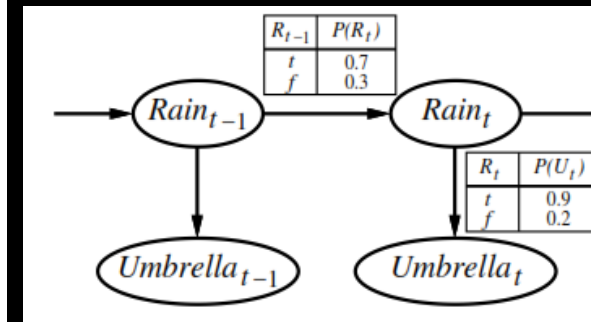
e.g., with  $U_1 = \text{true}$ ,  $O_1 = \begin{pmatrix} 0.9 & 0 \\ 0 & 0.2 \end{pmatrix}$

Forward and backward messages as column vectors:

$$\mathbf{f}_{1:t+1} = \alpha \mathbf{O}_{t+1} \mathbf{T}^\top \mathbf{f}_{1:t}$$


$$\mathbf{b}_{k+1:t} = \mathbf{T} \mathbf{O}_{k+1} \mathbf{b}_{k+2:t}$$

Forward-backward algorithm needs time  $O(S^2t)$  and space  $O(St)$



# HMMs

- Since we only have one variable (the system state), HMMs are often depicted like state machines:



Don't confuse this with a Bayes net diagram!

HMM with four states: A, B, C, D

- In robotics: HMMs useful for anomaly detection, localization, etc.



# Example: The Dishonest Casino

A casino has two dice:

- Fair die

$$P(1) = P(2) = P(3) = P(5) = P(6) = 1/6$$

- Loaded die

$$P(1) = P(2) = P(3) = P(5) = 1/10$$

$$P(6) = 1/2$$



Casino player switches back-&-forth between fair and loaded die once every 20 turns

## Game:

1. You bet \$1
2. You roll (always with a fair die)
3. Casino player rolls (maybe with fair die, maybe with loaded die)
4. Highest number wins \$2

# Question # 1 – Evaluation

## GIVEN

A sequence of rolls by the casino player:

1245526462146146136661664661636616366163616515615115146123562344



$$\text{Prob} = 1.3 \times 10^{-35}$$

## QUESTION

How likely is this sequence, given our model of how the casino works?

This is the **EVALUATION** problem in HMMs

## Question # 2 – Decoding

### GIVEN

A sequence of rolls by the casino player

1	2	4	5	5	2	6	4	6	2	1	4	6	1	4	6	1	3	6	1	3	6	6	6	1	6	6	4	6	6	1	6	3	6	6	1	6	3	6	6	1	6	3	6	1	6	5	1	5	6	1	5	1	1	5	1	4	6	1	2	3	5	6	2	3	4	4
FAIR															LOADED															FAIR																																				

### QUESTION

What portion of the sequence was generated with the fair die, and what portion with the loaded die?


This is the **DECODING** question in HMMs

# Question # 3 – Learning

## GIVEN

A sequence of rolls by the casino player

1245526462146146136136661664661636616366163616515615115146123562344


$$\text{Prob}(6) = 64\%$$

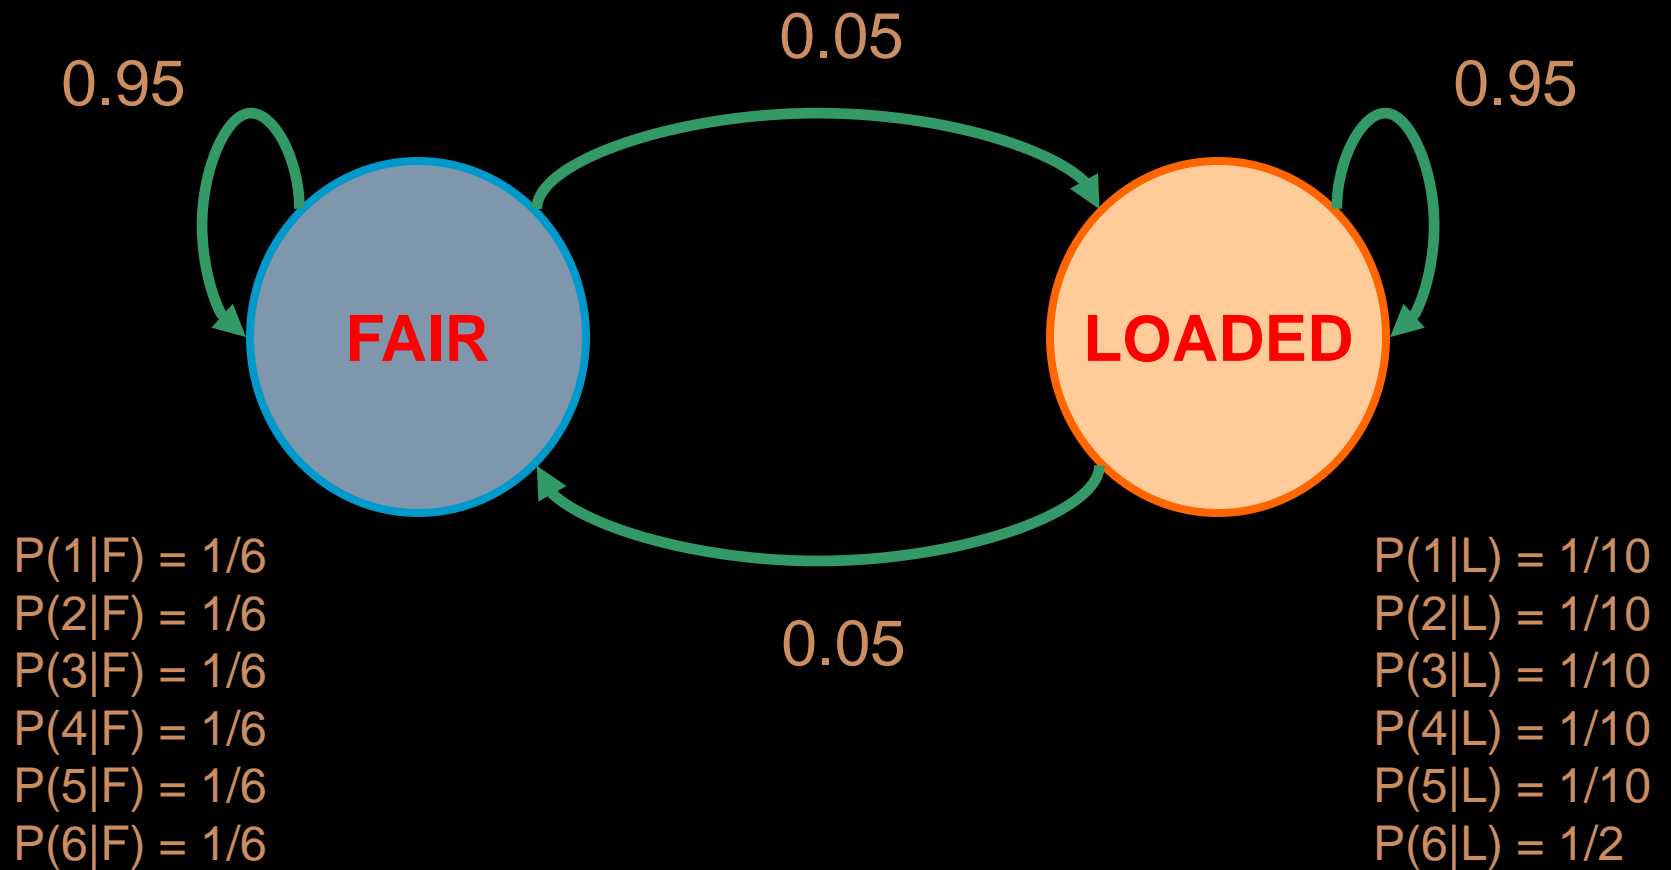
## QUESTION

How “loaded” is the loaded die? How “fair” is the fair die? How often does the casino player change from fair to loaded, and back?

This is the **LEARNING** question in HMMs


- We won't cover this, but if you're interested look-up the Baum–Welch Algorithm

# The dishonest casino HMM model



## Not just for casino games!

- We need methods like this for anomaly detection in robotics!



- We'll see an example later

# More precise definition of a hidden Markov model

**Definition:** A hidden Markov model (HMM)

- **Alphabet**  $\Sigma = \{ b_1, b_2, \dots, b_M \}$
- **Set of states**  $Q = \{ 1, \dots, K \}$
- **Transition probabilities** between any two states


$a_{ij}$  = transition prob from state  $i$  to state  $j$

$a_{i1} + \dots + a_{iK} = 1$ , for all states  $i = 1 \dots K$

- **Start probabilities**  $a_{0i}$   
 $a_{01} + \dots + a_{0K} = 1$
- **Emission probabilities** within each state

$e_i(b) = P(E_i = b \mid X_i = k)$


$e_i(b_1) + \dots + e_i(b_M) = 1$ , for all states  $i = 1 \dots K$



# A parse of a sequence

Given a sequence  $\mathbf{e} = e_1, \dots, e_N$ ,

A parse of  $\mathbf{e}$  is a sequence of states  $\mathbf{X} = x_1, \dots, x_N$





# Evaluation: Likelihood of a parse

- Given a sequence  $\mathbf{e} = e_1 \dots e_N$  and a parse  $\mathbf{x} = x_1, \dots, x_N$ ,
- To find how likely this scenario is (given our HMM):

$$P(\mathbf{e}, \mathbf{x}) = P(e_1, \dots, e_N, x_1, \dots, x_N) =$$

$$P(e_N | x_N) P(x_N | x_{N-1}) \dots P(e_2 | x_2) P(x_2 | x_1) P(e_1 | x_1) P(x_1) =$$

$$a_{0x_1} a_{x_1x_2} \dots a_{x_{N-1}x_N} e_{x_1}(e_1) \dots e_{x_N}(e_N)$$

# Evaluation problem for the dishonest casino

- Let the sequence of rolls be:

$$e = 1, 2, 1, 5, 6, 2, 1, 5, 2, 4$$



- Then, what is the likelihood of

**x** = Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair, Fair?

(say initial probs  $a_{0\text{Fair}} = \frac{1}{2}$ ,  $a_{0\text{Loaded}} = \frac{1}{2}$ )

$$\frac{1}{2} \times P(1 \mid \text{Fair}) P(\text{Fair} \mid \text{Fair}) P(2 \mid \text{Fair}) P(\text{Fair} \mid \text{Fair}) \dots P(4 \mid \text{Fair}) =$$

$$\frac{1}{2} \times (1/6)^{10} \times (0.95)^9 = .00000000521158647211 \sim 0.5 \times 10^{-9}$$

# Evaluation problem for the dishonest casino

So, the likelihood the die is fair in this run  
is just  $0.521 \times 10^{-9}$



What is the likelihood of

**x** = Loaded, Loaded, Loaded, Loaded, Loaded, Loaded, Loaded,  
Loaded, Loaded, Loaded?

$\frac{1}{2} \times P(1 \mid \text{Loaded}) P(\text{Loaded, Loaded}) \dots P(4 \mid \text{Loaded}) =$

$\frac{1}{2} \times (1/10)^9 \times (1/2)^1 (0.95)^9 = .00000000015756235243 \approx 0.16 \times 10^{-9}$

Therefore, it is more likely that all the rolls are done with the fair die,  
than that they are all done with the loaded die

# Evaluation problem for the dishonest casino

- Let the sequence of rolls be:

$$e = 1, 6, 6, 5, 6, 2, 6, 6, 3, 6$$



- What is the likelihood  $\mathbf{x} = F, F, \dots, F$ ?

$$\frac{1}{2} \times (1/6)^{10} \times (0.95)^9 \approx 0.5 \times 10^{-9}, \text{ same as before}$$


- What is the likelihood  $\mathbf{x} = L, L, \dots, L$ ?

$$\frac{1}{2} \times (1/10)^4 \times (1/2)^6 (0.95)^9 = .00000049238235134735 \approx 0.5 \times 10^{-7}$$

So, it is 100 times more likely the die is loaded

# Decoding Problem


- GIVEN  $\mathbf{e} = e_1, e_2, \dots, e_N$
- Find  $\mathbf{x} = x_1, \dots, x_N$ , that maximizes  $P(\mathbf{e}, \mathbf{x})$
- $x^* = \operatorname{argmax}_{\mathbf{x}} P(\mathbf{e}, \mathbf{x})$
- This is just like finding the max-likelihood sequence of a Bayes net
- Use the Viterbi algorithm



BREAK

# HMMs for Robot Localization


- What we need to know



- (1) State domains: e.g. (x,y) position
- (2) Evidence domains: Sensor model
- (3) Probability of states at time 0: Usually known or “kidnapped robot”
- (4) Transition probability: Defined by actions
- (5) Emission probability: Defined by sensor model

# Localization Example


*Example from  
Michael Pfeiffer*




Sensor model: never more than 1 mistake


Motion model: may not execute action with small prob.







Sometimes  
sensors are  
wrong




Sometimes  
sensors are  
wrong

Sometimes we  
have actuation  
error

# Wall-sensing Localization Example

*Example from  
Michael Pfeiffer*




$t=0$

Sensor model: never more than 1 mistake


Motion model: may not execute action with small prob.

# Wall-sensing Localization Example




$t=1$

# Wall-sensing Localization Example




t=2

# Wall-sensing Localization Example



$t=3$


# Wall-sensing Localization Example



$t=4$



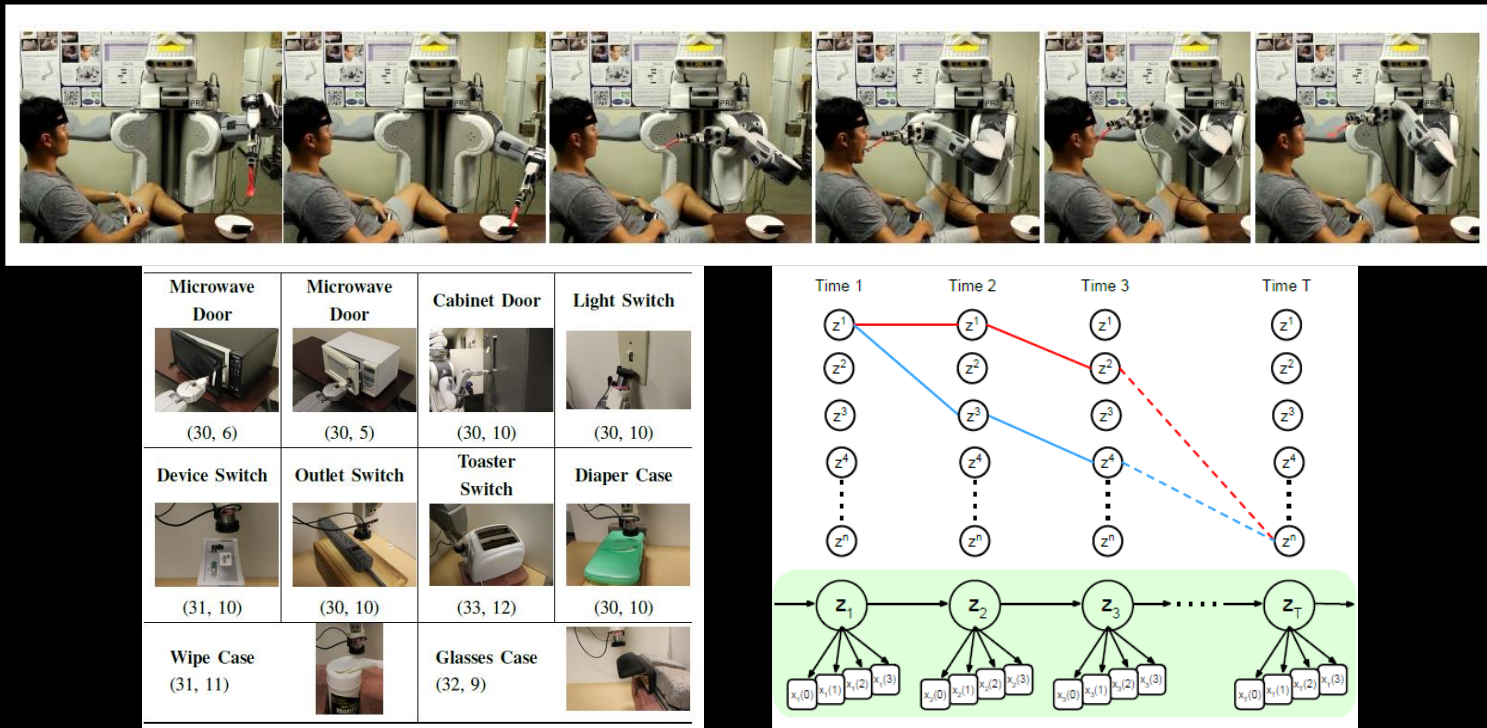
# Wall-sensing Localization Example



$t=5$

# HMM Evaluation Problem Application: Anomaly Detection for Robot Manipulation

- Monitor anomalies in manipulation task execution with an HMM
  - Uses force, visual, auditory, and kinematic sensing



- Using General Hidden Markov Model library (<http://www.ghmm.org>)

“Multimodal Execution Monitoring for Anomaly Detection During Robot Manipulation,”  
 D. Park, Z. Erickson, T. Bhattacharjee, and C. Kemp, ICRA 2016

# Summary

- Can construct Bayes nets that account for time-varying uncertainty
- These are useful for filtering, prediction, smoothing and most likely path questions
- HMMs are Bayes nets that make Markov assumption and have a single discrete random variable
  - Evaluation: How likely is this sequence of evidence given a model
  - Decoding: Find sequence of states that best explains the evidence
- HMMs have been used for localization and anomaly detection in robotics

# Homework

- AI book Ch. 15.4-15.5
- Kalman Filter Derivation [here](#)
- HW 4 is due Wednesday!